

Naive Bayes (Gaussian) Credit Card Fraud Detecting Project

Nawal Obaid

9/12/2022

PROBLEM STATEMENT

- Credit card companies need to have the ability to recognize fraudulent credit card transactions so that customers are not charged for items that they did not purchase.
- Datasets contains transactions made by credit cards in September 2013 by european cardholders. This dataset presents transactions that occurred in two days, where we have 492 frauds out of 284,807 transactions. The dataset is highly unbalanced, the positive class (frauds) account for 0.172% of all transactions.
- The data contains only numerical input variables which are the result of a PCA transformation. Unfortunately, due to confidentiality issues, we cannot provide the original features and more background information about the data.
- Input Features: V1, V2, ... V28 are the principal components obtained with PCA, the only features which have not been transformed with PCA are 'Time' and 'Amount'. Feature 'Time' contains the seconds elapsed between each transaction and the first transaction in the dataset. The feature 'Amount' is the transaction Amount, this feature can be used for example-dependant cost-sensitive learning.
- Output: 1 in case of fraud and 0 otherwise.
- Link to the dataset: <https://www.kaggle.com/mlg-ulb/creditcardfraud/home>

Import the needed libraries

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
```

Get the data

```
In [2]: card = pd.read_csv('creditcard.csv')
```

```
In [3]: card.head()
```

```
Out[3]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	...	-0.01

1	0.0	1.191857	0.266151	0.166480	0.448154	-0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.22
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.24
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	-0.10
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	-0.00

5 rows × 31 columns

In [4]: `card.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   Time        284807 non-null float64
 1   V1          284807 non-null float64
 2   V2          284807 non-null float64
 3   V3          284807 non-null float64
 4   V4          284807 non-null float64
 5   V5          284807 non-null float64
 6   V6          284807 non-null float64
 7   V7          284807 non-null float64
 8   V8          284807 non-null float64
 9   V9          284807 non-null float64
10  V10         284807 non-null float64
11  V11         284807 non-null float64
12  V12         284807 non-null float64
13  V13         284807 non-null float64
14  V14         284807 non-null float64
15  V15         284807 non-null float64
16  V16         284807 non-null float64
17  V17         284807 non-null float64
18  V18         284807 non-null float64
19  V19         284807 non-null float64
20  V20         284807 non-null float64
21  V21         284807 non-null float64
22  V22         284807 non-null float64
23  V23         284807 non-null float64
24  V24         284807 non-null float64
25  V25         284807 non-null float64
26  V26         284807 non-null float64
27  V27         284807 non-null float64
28  V28         284807 non-null float64
29  Amount      284807 non-null float64
30  Class       284807 non-null int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB
```

In [5]: `card.describe()`


	Time	V1	V2	V3	V4	V5	V6
count	284807.000000	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	94813.859575	3.918649e-15	5.682686e-16	-8.761736e-15	2.811118e-15	-1.552103e-15	2.040130e-15
std	47488.145955	1.958696e+00	1.651309e+00	1.516255e+00	1.415869e+00	1.380247e+00	1.332271e+00
min	0.000000	-5.640751e+01	-7.271573e+01	-4.832559e+01	-5.683171e+00	-1.137433e+02	-2.616051e+01
25%	54201.500000	-9.203734e-01	-5.985499e-01	-8.903648e-01	-8.486401e-01	-6.915971e-01	-7.682956e-01
50%	84692.000000	1.810880e-02	6.548556e-02	1.798463e-01	-1.984653e-02	-5.433583e-02	-2.741871e-01

75%	139320.500000	1.315642e+00	8.037239e-01	1.027196e+00	7.433413e-01	6.119264e-01	3.985649e-01
max	172792.000000	2.454930e+00	2.205773e+01	9.382558e+00	1.687534e+01	3.480167e+01	7.330163e+01

8 rows × 31 columns

```
In [6]: sns.heatmap(card.isnull())
# no missing vaules
```

```
Out[6]: <AxesSubplot:>
```



```
In [7]: is_fraud = card[card['Class']==1]
not_fraud = card[card['Class']==0]
```

```
In [8]: len(is_fraud)
```

```
Out[8]: 492
```

```
In [9]: len(not_fraud)
# imbalanced dataset
```


```
Out[9]: 284315
```

```
In [10]: # let's get counts and percentages!
print('Number of spam emails:', len(is_fraud))
print('Percentage of spam emails:', len(is_fraud)/len(card) * 100, '%')
print('Number of not spam emails:', len(not_fraud))
print('Percentage of not spam emails:', len(not_fraud)/len(card) * 100, '%')
```

```
Number of spam emails: 492
Percentage of spam emails: 0.1727485630620034 %
Number of not spam emails: 284315
Percentage of not spam emails: 99.82725143693798 %
```

```
In [11]: c = sns.countplot(data = card, x = card['Class'])
# spam is 1, not-spam or ham is 0
c.set_xlabel('Fraud vs. Not Fraud')
c.set_ylabel('Counts')
```


```
Out[11]: Text(0, 0.5, 'Counts')
```



```
In [12]: corr = card.corr()
```

```
In [13]: plt.figure(figsize = (30,20))
sns.heatmap(corr, annot = True)
# it seems there is no correlection among columns in this dataset since these columns ar
```

```
Out[13]: <AxesSubplot:>
```



Feature Engineering

```
In [14]: # Feature scale/normalize the "Amount" column
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
card['Amount_Norm'] = sc.fit_transform(card['Amount'].values.reshape(-1,1))
```

```
In [15]: card.head()
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	...	0.27
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.63
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.77
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	0.00
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	0.79

5 rows × 32 columns

```
In [16]: # we do not need the 'Amount' column anymore so we can drop it
card.drop('Amount' , axis = 1, inplace = True)
```

```
In [17]: card.head()
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	...	-0.01
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.22
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.24
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	-0.10
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	-0.00

5 rows × 31 columns

Train Test Split

```
In [18]: X = card.drop('Class', axis = 1)
y = card['Class']
```

```
In [19]: from sklearn.model_selection import train_test_split
```

```
In [20]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

```
In [21]: X_train.shape
```

```
Out[21]: (227845, 30)
```

```
In [22]: X_test.shape
```

```
Out[22]: (56962, 30)
```

```
In [23]: y_train.shape
```

```
Out[23]: (227845,)
```

```
In [24]: y_test.shape
```

```
Out[24]: (56962,)
```

```
In [25]: from sklearn.naive_bayes import GaussianNB
classifier = GaussianNB()
```

```
In [26]: classifier.fit(X_train, y_train)
```

```
Out[26]: ▼ GaussianNB
GaussianNB()
```

```
In [27]: predictions = classifier.predict(X_test)
```

```
In [28]: from sklearn.metrics import classification_report, confusion_matrix
```

```
In [54]: print('*****Classification Report*****')
print(classification_report(y_test, predictions))
print('*****Confusion Matrix*****')
print(confusion_matrix(y_test, predictions))
```

```
*****Classification Report*****
              precision    recall  f1-score   support


     0           1.00       0.98       0.99       56871
     1           0.08       0.87       0.15         91

 accuracy          0.98       0.98       0.98       56962
 macro avg          0.54       0.93       0.57       56962
 weighted avg       1.00       0.98       0.99       56962

*****Confusion Matrix*****
[[55997   874]
 [    12    79]]
```

```
In [55]: cm = confusion_matrix(y_test, predictions)
sns.heatmap(cm, annot = True, fmt = "d", cmap = 'Reds')
```

```
Out[55]: <AxesSubplot:>
```



The model misclassified about 401 cases (32 of Error Type II), a big reason is that the dataset is originally imbalanced!!

Improve the model

```
In [31]: # we can drop the features we are not interested in anymore since they do not make a gre

X = card.drop(['Class','Time','V8', 'V13', 'V15', 'V20', 'V22', 'V23', 'V24','V25', 'V26
y = card['Class']
```

```
In [40]: # train and test the model again based on the new X
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

```
In [45]: GNBclassifier = GaussianNB()
GNBclassifier.fit(X_train, y_train)
predictions = GNBclassifier.predict(X_test)
```

```
In [48]: print('*****Classification Report*****')
print(classification_report(y_test, predictions))
print('*****Confusion Matrix*****')
print(confusion_matrix(y_test, predictions))
```

```
*****Classification Report*****
              precision    recall  f1-score   support

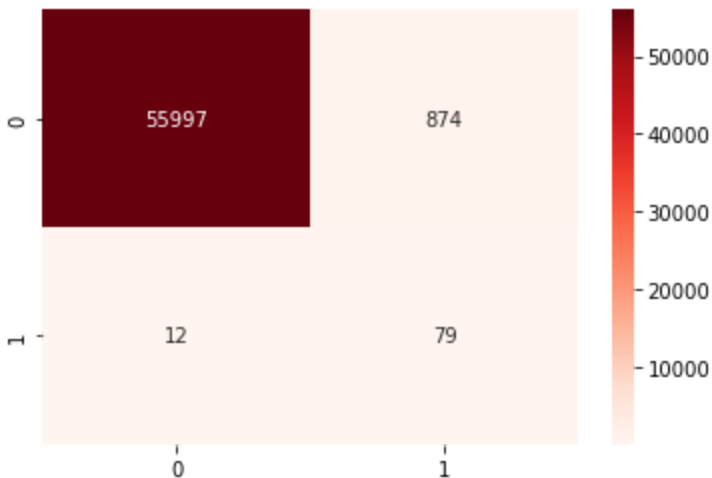
     0           1.00       0.98       0.99       56871
     1           0.08       0.87       0.15         91

 accuracy          0.98       0.98       0.98       56962
 macro avg          0.54       0.93       0.57       56962
 weighted avg          1.00       0.98       0.99       56962

*****Confusion Matrix*****
[[55997   874]
 [    12    79]]
```

```
In [56]: cm = confusion_matrix(y_test, predictions)
sns.heatmap(cm, annot = True, fmt = "d", cmap = 'Reds')
# The model misclassified 12 cases of Type II which is kind of improvement
```

```
Out[56]: <AxesSubplot:>
```



```
In [62]: print('Number of fraud cases/points in the testing dataset = ', sum(y_test), ' the model

Number of fraud cases/points in the testing dataset = 91 the model has misclassified 1
2 cases
```