

# Support Vector Machines - Bank Customers Retirement Predictions Project

Nawal Obaid

9/7/2022

## PROBLEM STATEMENT

You work as a data scientist at a major bank in NYC and you have been tasked to develop a model that can predict whether a customer is able to retire or not based on his/her features. Features are his/her age and net 401K savings (retirement savings in the U.S.). You thought that Support Vector Machines can be a great candidate to solve the problem.

## Import Needed Libraries

```
In [97]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
```

## Load Data

```
In [98]: bank_df = pd.read_csv('Bank_Customer_retirement.csv')
```

```
In [99]: bank_df.head()
# need to do scaling to this dataset
# need to use SVM as a classifier
```

```
Out[99]:
```

	Customer ID	Age	401K Savings	Retire
0	0	39.180417	322349.8740	0
1	1	56.101686	768671.5740	1
2	2	57.023043	821505.4718	1
3	3	43.711358	494187.4850	0
4	4	54.728823	691435.7723	1

```
In [100]: bank_df.info()
#500 entries, (total 4 columns)

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 4 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Customer ID     500 non-null   int64
1   Age             500 non-null   float64
2   401K Savings    500 non-null   float64
3   Retire          500 non-null   int64
```

```
dtypes: float64(2), int64(2)  
memory usage: 15.8 KB
```

```
In [101... bank_df.shape
```

```
Out[101]: (500, 4)
```

```
In [102... bank_df.describe()
```


```
Out[102]:
```

	Customer ID	Age	401K Savings	Retire
count	500.000000	500.000000	500.000000	500.000000
mean	249.500000	46.757077	534984.465804	0.500000
std	144.481833	7.790125	187675.818881	0.500501
min	0.000000	25.000000	10.000000	0.000000
25%	124.750000	41.299451	382626.524425	0.000000
50%	249.500000	46.695770	534512.984350	0.500000
75%	374.250000	52.322551	680670.257025	1.000000
max	499.000000	70.000000	1000000.000000	1.000000

## EDA


```
In [103... sns.pairplot(bank_df, hue = 'Retire', vars = ['Age', '401K Savings'], palette = 'plasma')
```

```
Out[103]: <seaborn.axisgrid.PairGrid at 0x272d20b7910>
```




```
In [104... sns.scatterplot(data= bank_df, x='Age', y='401K Savings', hue = 'Retire', palette = 'pla
```

```
Out[104]: <AxesSubplot:xlabel='Age', ylabel='401K Savings'>
```



```
In [105]: plt.figure(figsize=[8, 8])
sns.countplot( data = bank_df, x='Retire', palette='plasma')
```

```
Out[105]: <AxesSubplot:xlabel='Retire', ylabel='count'>
```



```
In [106]: # Let's explore which dataset is missing
sns.heatmap(bank_df.isnull(), yticklabels = False, cbar = False, cmap="Blues")
# No missing values
# No data imputation needed
```

```
Out[106]: <AxesSubplot:>
```



Customer ID      Age      401K Savings      Retire

```
In [107... # Let's drop the Id coloumn
bank_df.drop('Customer ID',axis=1,inplace=True)
```

```
In [108... bank_df.head(2)
```

```
Out[108]:
```

	Age	401K Savings	Retire
0	39.180417	322349.874	0
1	56.101686	768671.574	1

```
In [109... X = bank_df.drop('Retire', axis =1)
y = bank_df['Retire']
```

## Feature Scaling/Normalization

```
In [110... # Scaling is needed
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X = scaler.fit_transform(X)
```

## Train Test Split

```
In [111... from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=5)
```

## Tain/Fit the model

```
In [112... from sklearn.svm import SVC
```

```
In [113... classifier = SVC()
```

```
In [114... classifier.fit(X_train, y_train)
```

```
Out[114]:
```

▼ SVC

SVC()

```
In [115... predictions = classifier.predict(X_test)
```

```
In [116... from sklearn.metrics import classification_report, confusion_matrix
```

```
In [117... print('*****Classification Report*****')
print(classification_report(y_test, predictions))
print('*****Confusion Matrix*****')
print(confusion_matrix(y_test, predictions))
```

```
*****Classification Report*****
              precision    recall  f1-score   support

     0       0.95       0.93       0.94         43
     1       0.95       0.96       0.96         57

   accuracy              0.95         100
  macro avg              0.95         100
weighted avg              0.95         100

*****Confusion Matrix*****
[[40  3]
 [ 2 55]]
```


## The confusion matrix shows:

False Positives(Type I Error)= 3

False Negatives(Type II Error)= 2

```
In [119... cm = confusion_matrix(y_test, predictions)
sns.heatmap(cm, annot = True, fmt = "d", cmap = 'Blues')
```

Out[119]: <AxesSubplot:>



In [ ]: